

Machine Learning and Data Mining COMP 5318

Multi-class Classification


Classification Based on Probability

- Instead of just predicting the class, give the probability of the instance being that class
 - i.e., learn $p(y | x)$

- Recall that:

$$0 \leq p(\text{event}) \leq 1$$

$$p(\text{event}) + p(\neg \text{event}) = 1$$



Logistic Regression

- Takes a probabilistic approach to learning discriminative functions (i.e., a classifier)

- $h_{\theta}(x)$ should give $p(y = 1 | x; \theta)$

– Want $0 \leq h_{\theta}(x) \leq 1$


Can't just use linear regression with a threshold

- Logistic regression model:

$$h_{\theta}(x) = \sigma(\theta^T x)$$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$



Interpretation of Hypothesis Output

$$h_{\theta}(x) = \text{estimated } p(y = 1 | x; \theta)$$

Example: Cancer diagnosis from tumor size

$$x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} \quad \sigma \begin{bmatrix} 1 \\ \text{tumorSize} \end{bmatrix}$$

$$h_{\theta}(x) = 0.7$$

70% chance of tumor being malignant

Note that: $p(y = 0 | x; \theta) + p(y = 1 | x; \theta) = 1$

Therefore, $p(y = 0 | x; \theta) = 1 - p(y = 1 | x; \theta)$

Another Interpretation

- Equivalently, logistic regression assumes that

$$\log \frac{p(y=1 | \mathbf{x}; \boldsymbol{\theta})}{p(y=0 | \mathbf{x}; \boldsymbol{\theta})} = \theta_0 + \theta_1 x_1 + \dots + \theta_d x_d$$

if $y=1$

Side Note The odds in favor of an event is the quantity $\frac{p}{1-p}$, where p is the probability of the event. Example: If I toss a fair dice, what are the odds that I will have a 6?

- In other words, logistic regression assumes that the log odds is a linear function of \mathbf{x}


Based on slide by XiaoLi Fern

5

Logistic Regression

$$h_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$$


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



$\boldsymbol{\theta}^T \mathbf{x}$ should be large negative values for negative instances

$\boldsymbol{\theta}^T \mathbf{x}$ should be large positive values for positive instances

- Assume a threshold and...
 - Predict $y=1$ if $h_{\boldsymbol{\theta}}(\mathbf{x}) \geq 0.5$
 - Predict $y=0$ if $h_{\boldsymbol{\theta}}(\mathbf{x}) < 0.5$




Based on slide by Andrew Ng

6

Non-Linear Decision Boundary

- Can apply basis function expansion to features as with linear regression

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} \rightarrow \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_1 x_2 \\ x_1^2 \\ x_2^2 \\ x_1^2 x_2 \\ x_1 x_2^2 \\ \vdots \end{bmatrix}$$



7

Logistic Regression

- Model: $h_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$
- $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)} \in \mathbb{R}^d, y^{(i)} \in \{0, 1\}$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\boldsymbol{\theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix}$$

$$\mathbf{x}^T = \begin{bmatrix} 1 & x_1 & \dots & x_d \end{bmatrix}$$

8

Logistic Regression Objective Function

- Can't just use squared loss as in linear regression:

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n \left(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right)^2$$

- Using the logistic regression model

$$h_{\theta}(\mathbf{x}) = \frac{1}{1 + e^{-\theta^T \mathbf{x}}}$$

results in a non-convex optimization

9

Deriving the Cost Function via Maximum Likelihood Estimation

- Likelihood of data is given by: $l(\theta) = \prod_{i=1}^n p(y^{(i)} | \mathbf{x}^{(i)}; \theta)$

- So, looking for the θ that maximizes the likelihood

$$\theta_{\text{MLE}} = \arg \max_{\theta} l(\theta) = \arg \max_{\theta} \prod_{i=1}^n p(y^{(i)} | \mathbf{x}^{(i)}; \theta)$$

- Can take the log without changing the solution:

$$\begin{aligned} \theta_{\text{MLE}} &= \arg \max_{\theta} \log \prod_{i=1}^n p(y^{(i)} | \mathbf{x}^{(i)}; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^n \log p(y^{(i)} | \mathbf{x}^{(i)}; \theta) \end{aligned}$$

10

Deriving the Cost Function via Maximum Likelihood Estimation

- Expand as follows:

$$\begin{aligned} \theta_{\text{MLE}} &= \arg \max_{\theta} \sum_{i=1}^n \log p(y^{(i)} | \mathbf{x}^{(i)}; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^n \left[y^{(i)} \log p(y^{(i)} = 1 | \mathbf{x}^{(i)}; \theta) + (1 - y^{(i)}) \log (1 - p(y^{(i)} = 1 | \mathbf{x}^{(i)}; \theta)) \right] \end{aligned}$$

- Substitute in model, and take negative to yield

Logistic regression objective:

$$\min_{\theta} J(\theta)$$

$$J(\theta) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\theta}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(\mathbf{x}^{(i)})) \right]$$

11

Intuition Behind the Objective

$$J(\theta) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\theta}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(\mathbf{x}^{(i)})) \right]$$

- Cost of a single instance:

$$\text{cost}(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

- Can re-write objective function as


$$J(\theta) = \sum_{i=1}^n \text{cost}(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)})$$

Compare to linear regression: $J(\theta) = \frac{1}{2n} \sum_{i=1}^n \left(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right)^2$

12

Intuition Behind the Objective


$$\text{cost}(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



13

Intuition Behind the Objective

$$\text{cost}(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



If $y = 1$

- Cost = 0 if prediction is correct
- As $h_{\theta}(\mathbf{x}) \rightarrow 0$, cost $\rightarrow \infty$
- Captures intuition that larger mistakes should get larger penalties
 - e.g., predict $h_{\theta}(\mathbf{x}) = 0$, but $y = 1$

Based on example by Andrew Ng


14

Intuition Behind the Objective

$$\text{cost}(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

If $y = 0$

- Cost = 0 if prediction is correct
- As $(1 - h_{\theta}(\mathbf{x})) \rightarrow 0$, cost $\rightarrow \infty$
- Captures intuition that larger mistakes should get larger penalties



Based on example by Andrew Ng

15

Regularized Logistic Regression

$$J(\theta) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\theta}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(\mathbf{x}^{(i)})) \right]$$

- We can regularize logistic regression exactly as before:

$$\begin{aligned} J_{\text{regularized}}(\theta) &= J(\theta) + \frac{\lambda}{2} \sum_{j=1}^d \theta_j^2 \\ &= J(\theta) + \frac{\lambda}{2} \|\theta_{[1:d]}\|_2^2 \end{aligned}$$

16

Gradient Descent for Logistic Regression

$$J_{\text{reg}}(\theta) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\theta}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(\mathbf{x}^{(i)})) \right] + \frac{\lambda}{2} \|\theta_{[1:d]}\|_2^2$$

Want $\min_{\theta} J(\theta)$

- Initialize θ
- Repeat until convergence

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad \text{simultaneous update for } j = 0 \dots d$$

Use the natural logarithm ($\ln = \log_e$) to cancel with the $\exp()$ in $h_{\theta}(\mathbf{x})$

17

Gradient Descent for Logistic Regression

$$J_{\text{reg}}(\theta) = - \sum_{i=1}^n \left[y^{(i)} \log h_{\theta}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(\mathbf{x}^{(i)})) \right] + \frac{\lambda}{2} \|\theta_{[1:d]}\|_2^2$$

Want $\min_{\theta} J(\theta)$

- Initialize θ
- Repeat until convergence (simultaneous update for $j = 0 \dots d$)

$$\begin{aligned} \theta_0 &\leftarrow \theta_0 - \alpha \sum_{i=1}^n \left(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right) \\ \theta_j &\leftarrow \theta_j - \alpha \left[\sum_{i=1}^n \left(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right) x_j^{(i)} + \lambda \theta_j \right] \end{aligned}$$

18

Gradient Descent for Logistic Regression

- Initialize θ
- Repeat until convergence (simultaneous update for $j = 0 \dots d$)

$$\begin{aligned} \theta_0 &\leftarrow \theta_0 - \alpha \sum_{i=1}^n \left(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right) \\ \theta_j &\leftarrow \theta_j - \alpha \left[\sum_{i=1}^n \left(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right) x_j^{(i)} + \lambda \theta_j \right] \end{aligned}$$

This looks IDENTICAL to linear regression!!!


- Ignoring the $1/n$ constant
- However, the form of the model is very different:

$$h_{\theta}(\mathbf{x}) = \frac{1}{1 + e^{-\theta^T \mathbf{x}}}$$


19

Multi-Class Classification

Binary classification:



Multi-class classification:



Disease diagnosis: healthy / cold / flu / pneumonia

Object classification: desk / chair / monitor / bookcase

20

Multi-Class Logistic Regression

- For 2 classes:

$$h_{\theta}(\mathbf{x}) = \frac{1}{1 + \exp(-\theta^T \mathbf{x})} = \frac{\exp(\theta^T \mathbf{x})}{\boxed{1} + \exp(\theta^T \mathbf{x})}$$

weight assigned to $y = 0$
weight assigned to $y = 1$

- For C classes $\{1, \dots, C\}$:


$$p(y = c \mid \mathbf{x}; \theta_1, \dots, \theta_C) = \frac{\exp(\theta_c^T \mathbf{x})}{\sum_{c=1}^C \exp(\theta_c^T \mathbf{x})}$$

- Called the **softmax** function

21

Multi-Class Logistic Regression

Split into One vs Rest:



- Train a logistic regression classifier for each class i to predict the probability that $y = i$ with

$$h_c(\mathbf{x}) = \frac{\exp(\theta_c^T \mathbf{x})}{\sum_{c=1}^C \exp(\theta_c^T \mathbf{x})}$$

22

Implementing Multi-Class Logistic Regression

- Use $h_c(\mathbf{x}) = \frac{\exp(\theta_c^T \mathbf{x})}{\sum_{c=1}^C \exp(\theta_c^T \mathbf{x})}$ as the model for class c
- Gradient descent simultaneously updates all parameters for all models
 - Same derivative as before, just with the above $h_c(\mathbf{x})$
- Predict class label as the most probable label

$$\max_c h_c(\mathbf{x})$$

23